

Probability Concepts for Machine Learning Engineers

A practical academic tutorial on uncertainty, estimation, inference, and probabilistic decision-making.

Mohamad Zbib

July 3, 2026

How to read this note

Probability is not a bag of formulas. It is a language for disciplined uncertainty. In machine learning, we use it to describe data generation, quantify model confidence, train with noisy observations, compare hypotheses, and decide what action is worth taking.

Contents

1	The Core View	2
2	Events, Random Variables, and Distributions	2
2.1	Joint, Marginal, and Conditional Distributions	2
3	Expectation, Variance, and Covariance	2
4	Independence and Conditional Independence	3
5	Bayes' Rule	3
5.1	MAP and Maximum Likelihood	4
6	Likelihood, Log-Likelihood, and Cross-Entropy	4
7	Common Distributions	4
8	Sampling and Monte Carlo Thinking	5
9	The Law of Large Numbers and the Central Limit Theorem	5
10	Bias, Variance, and Estimation	5
11	Information, Entropy, and KL Divergence	5
12	Calibration and Uncertainty	6
13	Probabilistic Graphical Thinking	6
14	Decision Theory	6
15	Distribution Shift	6
16	A Compact Workflow	7
17	Formula Sheet	7
18	Closing Perspective	7

1 The Core View

Machine learning begins with a simple tension: the world is structured, but our observations are incomplete. Probability gives us a way to work with that tension without pretending uncertainty is noise to be ignored.

A probabilistic model says: there is an input space, an output space, and a mechanism that makes some outcomes more plausible than others. Learning means using data to improve that mechanism. Prediction means turning the mechanism into a statement about what may happen next.

The mental model

Think of probability as a scoring system over possible worlds. A model does not merely output an answer; it assigns mass or density to alternatives. Good ML engineering keeps track of which uncertainty is reducible by more data, which uncertainty is inherent, and which uncertainty comes from model assumptions.

2 Events, Random Variables, and Distributions

An event is a set of outcomes. A random variable is a function that maps an outcome to a value. A distribution tells us how probability is allocated across those values.

For a discrete random variable X , the probability mass function is

$$p_X(x) = \mathbb{P}(X = x).$$

For a continuous random variable, the probability density function $f_X(x)$ is not itself a probability; probabilities come from areas:

$$\mathbb{P}(a \leq X \leq b) = \int_a^b f_X(x) dx.$$

Engineering intuition

Discrete probability counts possibilities. Continuous probability measures regions. When working with losses, likelihoods, and gradients, forgetting this difference leads to many subtle bugs.

2.1 Joint, Marginal, and Conditional Distributions

The joint distribution $p(x, y)$ describes how variables behave together. The marginal distribution removes variables we are not currently asking about:

$$p(x) = \sum_y p(x, y) \quad \text{or} \quad p(x) = \int p(x, y) dy.$$

The conditional distribution focuses on a slice of the joint:

$$p(y | x) = \frac{p(x, y)}{p(x)}.$$

In supervised learning, the object we usually want is $p(y | x)$: the uncertainty about a target once features are known. In generative modeling, we often care about $p(x)$ or $p(x, y)$: the process that could have produced the data.

3 Expectation, Variance, and Covariance

Expectation is the long-run average value under a distribution:

$$\mathbb{E}[X] = \sum_x x p(x) \quad \text{or} \quad \mathbb{E}[X] = \int x f(x) dx.$$

More generally, for any function g ,

$$\mathbb{E}[g(X)] = \sum_x g(x)p(x) \quad \text{or} \quad \mathbb{E}[g(X)] = \int g(x)f(x) dx.$$

Variance measures spread:

$$\text{Var}(X) = \mathbb{E}[(X - \mathbb{E}[X])^2] = \mathbb{E}[X^2] - \mathbb{E}[X]^2.$$

Covariance measures how two variables move together:

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])].$$

Why ML engineers care

Loss functions are expectations. Training on mini-batches estimates those expectations. Generalization asks whether the empirical average is close to the population average. Optimization noise is often variance wearing a different name.

4 Independence and Conditional Independence

Two variables are independent when knowing one tells us nothing about the other:

$$p(x, y) = p(x)p(y).$$

They are conditionally independent given z when the relationship disappears after z is known:

$$p(x, y | z) = p(x | z)p(y | z).$$

Conditional independence is the heart of graphical models, Naive Bayes, hidden Markov models, Bayesian networks, and many scalable approximations. It lets us factor large distributions into smaller parts.

A useful warning

Independence is fragile. Variables can be marginally independent but conditionally dependent, or marginally dependent but conditionally independent. Conditioning changes the world you are talking about.

5 Bayes' Rule

Bayes' rule is the algebra of belief revision:

$$p(\theta | D) = \frac{p(D | \theta)p(\theta)}{p(D)}.$$

Each term has a role:

Term	Meaning
$p(\theta)$	Prior belief before seeing the data
$p(D \theta)$	Likelihood of the observed data under parameter θ
$p(D)$	Evidence or normalizing constant
$p(\theta D)$	Posterior belief after observing the data

For ML engineers, Bayes' rule is not only a theorem. It is a pattern: combine assumptions with evidence, then normalize.

5.1 MAP and Maximum Likelihood

Maximum likelihood chooses parameters that make the data most probable:

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} p(D \mid \theta).$$

Maximum a posteriori estimation includes a prior:

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} p(D \mid \theta)p(\theta).$$

Regularization often has a Bayesian reading. For example, an L_2 penalty corresponds to a Gaussian prior over weights; an L_1 penalty corresponds to a Laplace prior.

6 Likelihood, Log-Likelihood, and Cross-Entropy

If data points are assumed independent given θ , then

$$p(D \mid \theta) = \prod_{i=1}^n p(x_i \mid \theta).$$

Products of many probabilities become tiny, so we optimize the log-likelihood:

$$\log p(D \mid \theta) = \sum_{i=1}^n \log p(x_i \mid \theta).$$

Cross-entropy appears when the true distribution p is compared to a model distribution q :

$$H(p, q) = -\mathbb{E}_{x \sim p}[\log q(x)].$$

In classification, minimizing cross-entropy encourages the model to assign high probability to the correct label. It is not simply a classification penalty; it is a probabilistic scoring rule.

Practical takeaway

When a model outputs probabilities, the loss should reward calibrated confidence. Being right with low confidence and being wrong with high confidence are not equivalent mistakes.

7 Common Distributions

Distribution	Use it for	ML examples
Bernoulli	One binary event	Click/no click, fraud/not fraud, success/failure
Categorical	One of K classes	Multiclass labels, token selection
Binomial	Number of successes in fixed trials	A/B tests, repeated binary outcomes
Poisson	Counts over time or space	Arrivals, defects, rare events
Gaussian	Continuous values with symmetric noise	Regression residuals, latent variables, approximate posteriors
Exponential	Waiting times	Time-to-event modeling, reliability
Beta	Distribution over probabilities	Bayesian conversion rates, uncertainty over Bernoulli parameters
Dirichlet	Distribution over categorical probabilities	Topic models, class-proportion priors

8 Sampling and Monte Carlo Thinking

Many expectations cannot be computed exactly. Monte Carlo methods approximate them with samples:

$$\mathbb{E}[g(X)] \approx \frac{1}{m} \sum_{j=1}^m g(x^{(j)}), \quad x^{(j)} \sim p(x).$$

This idea appears everywhere: dropout, bootstrapping, Bayesian posterior approximation, policy gradients, simulation, uncertainty estimation, and offline evaluation.

The Monte Carlo tradeoff

More samples usually reduce variance but cost more compute. Better sampling strategies, such as importance sampling or stratified sampling, reduce wasted samples by focusing computation where it matters.

9 The Law of Large Numbers and the Central Limit Theorem

The law of large numbers says sample averages converge toward expectations:

$$\frac{1}{n} \sum_{i=1}^n X_i \rightarrow \mathbb{E}[X].$$

The central limit theorem says that, under broad conditions, normalized sums look approximately Gaussian:

$$\frac{\sum_{i=1}^n X_i - n\mu}{\sigma\sqrt{n}} \Rightarrow \mathcal{N}(0, 1).$$

These results justify confidence intervals, error bars, mini-batch estimates, and many approximate diagnostics. They also have limits: dependence, heavy tails, selection bias, and distribution shift can break the comfort they provide.

10 Bias, Variance, and Estimation

An estimator is a function of data. Its error can often be decomposed into bias and variance. Bias is systematic error; variance is sensitivity to the particular sample.

In model selection, a very flexible model may have low bias but high variance. A very restricted model may have high bias but low variance. Regularization, early stopping, data augmentation, ensembling, and better features are all ways of reshaping this tradeoff.

$$\mathbb{E}[(\hat{f}(x) - f(x))^2] = \text{Bias}[\hat{f}(x)]^2 + \text{Var}[\hat{f}(x)] + \sigma^2.$$

The final term σ^2 is irreducible noise: even the best model cannot predict what the data-generating process does not reveal.

11 Information, Entropy, and KL Divergence

Entropy measures uncertainty:

$$H(p) = - \sum_x p(x) \log p(x).$$

KL divergence measures the extra cost of using q when the true distribution is p :

$$D_{\text{KL}}(p||q) = \sum_x p(x) \log \frac{p(x)}{q(x)}.$$

KL is not symmetric, so $D_{\text{KL}}(p\|q)$ and $D_{\text{KL}}(q\|p)$ punish different mistakes. In variational inference and generative modeling, this asymmetry matters.

Interpretation

Entropy asks, “How uncertain is this distribution?” Cross-entropy asks, “How expensive is it to encode outcomes from p using q ?” KL divergence asks, “How much worse is q than the ideal code based on p ?”

12 Calibration and Uncertainty

A classifier is calibrated if its confidence matches reality. If it predicts 0.8 probability across many cases, about 80% of those cases should be correct.

Accuracy and calibration are different. A model can be accurate but overconfident, or less accurate but better calibrated. Calibration matters in medical triage, fraud review, ranking, autonomous systems, and any setting where the probability itself drives a decision.

Common tools include reliability diagrams, expected calibration error, temperature scaling, Bayesian neural networks, ensembles, conformal prediction, and abstention policies.

13 Probabilistic Graphical Thinking

Graphical models are useful because they make assumptions visible. A directed edge often means direct dependence. Missing edges often mean conditional independence.

For example, suppose user intent Z influences query text X and clicked item Y . A model may assume

$$p(x, y, z) = p(z)p(x | z)p(y | z).$$

This says text and click are dependent overall, but independent once intent is known. Whether that assumption is true is an engineering question, not just a mathematical one.

14 Decision Theory

Predictions become useful when connected to actions. Decision theory says choose the action with the best expected utility or lowest expected loss:

$$a^* = \arg \min_a \mathbb{E}[L(a, Y) | x].$$

This is why thresholds are business decisions, not just model metrics. A false positive and a false negative may have different costs. A probability estimate becomes valuable when it supports a cost-aware choice.

Thresholds

The default 0.5 threshold is rarely sacred. Tune thresholds against the real cost function, capacity constraints, legal requirements, and user experience.

15 Distribution Shift

Training data comes from one distribution; production data may come from another. This is distribution shift:

$$p_{\text{train}}(x, y) \neq p_{\text{prod}}(x, y).$$

Important forms include:

- Covariate shift: $p(x)$ changes while $p(y | x)$ is mostly stable.
- Label shift: $p(y)$ changes while $p(x | y)$ is mostly stable.
- Concept drift: $p(y | x)$ changes, often the hardest case.
- Selection bias: the observed data is not representative of the target population.

Monitoring should therefore track data features, labels when available, prediction distributions, calibration, delayed outcomes, and downstream decisions.

16 A Compact Workflow

1. Define the random variables: inputs, labels, latent factors, decisions, and costs.
2. State the modeling target: $p(y | x)$, $p(x)$, $p(x, y)$, a risk score, or an action.
3. Choose assumptions deliberately: independence, noise model, priors, stationarity, missingness.
4. Train by optimizing a probabilistic objective when probabilities matter.
5. Validate with both predictive metrics and uncertainty diagnostics.
6. Stress-test under shift, imbalance, missing data, and rare edge cases.
7. Connect probabilities to decisions using the real loss function.

17 Formula Sheet

Concept	Formula
Conditional probability	$p(a b) = p(a, b)/p(b)$
Product rule	$p(a, b) = p(a b)p(b) = p(b a)p(a)$
Marginalization	$p(a) = \sum_b p(a, b)$ or $p(a) = \int p(a, b) db$
Bayes' rule	$p(a b) = p(b a)p(a)/p(b)$
Expectation	$\mathbb{E}[X] = \sum_x xp(x)$ or $\int xf(x) dx$
Variance	$\text{Var}(X) = \mathbb{E}[X^2] - \mathbb{E}[X]^2$
Covariance	$\text{Cov}(X, Y) = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y]$
Entropy	$H(p) = -\sum_x p(x) \log p(x)$
Cross-entropy	$H(p, q) = -\mathbb{E}_p[\log q(X)]$
KL divergence	$D_{\text{KL}}(p q) = \mathbb{E}_p[\log p(X) - \log q(X)]$

18 Closing Perspective

Probability is the grammar behind machine learning systems that must operate under uncertainty. It helps us express what the model knows, what it assumes, what it has learned from data, and what remains uncertain.

The practical goal is not to make every system fully Bayesian or mathematically elaborate. The goal is to think clearly: define the variables, understand the distribution, optimize the right objective, validate uncertainty honestly, and make decisions with costs in view.

One sentence to remember

An ML model is not just a function from inputs to outputs; it is an engineered approximation to a probabilistic story about data, uncertainty, and decisions.